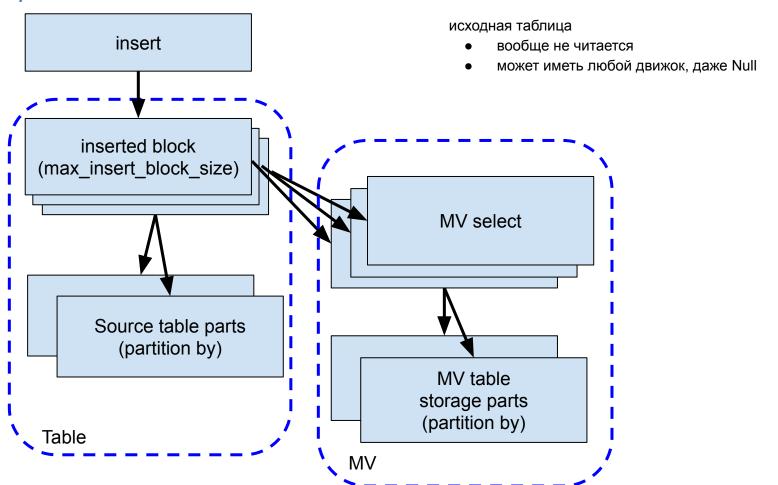
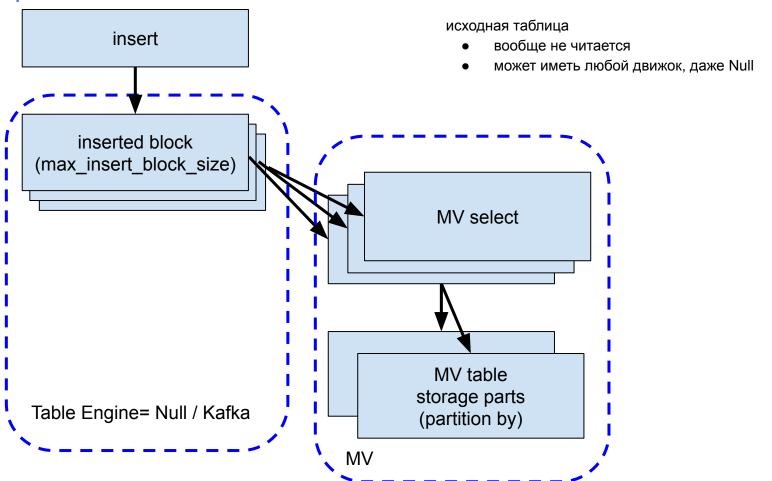


MV, на самом деле

MV изнутри



MV изнутри



Зачем MV нужны

- Пред-агрегации (Summing/Aggregating)
- Единая точка входа / Трансформации (ExtractTransform)
- Дублирование таблицы в таблицу с другим ключом (+эмуляция инверсного индекса)
- Kafka

Неявная таблица `.inner.mv1`

CREATE MATERIALIZED VIEW mv1

ENGINE = SummingMergeTree
PARTITION BY toYYYYMM(d)
ORDER BY (a, b, d)

AS

SELECT a, b, d, count() AS cnt FROM source

GROUP BY a, b, d;

DESCRIBE TABLE `.inner.mv1`

mame—	type	\neg default_type \neg
a	Int64	
b	Int64	
d	Date	
cnt	UInt64	

Явная таблица dest

CREATE TABLE dest

(a Int64, b Int64, d Date, cnt UInt64)

ENGINE = SummingMergeTree

PARTITION BY toYYYYMM(d)

ORDER BY (a, b, d);

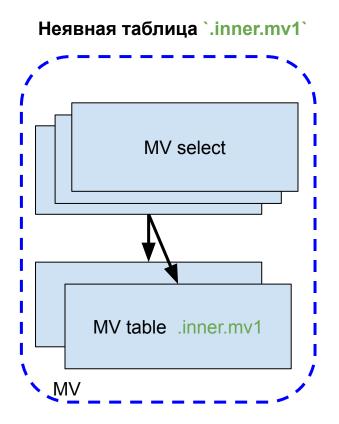
CREATE MATERIALIZED VIEW mv1

TO dest

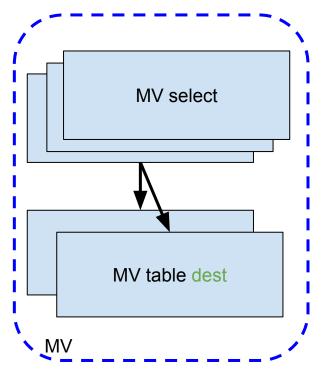
AS

SELECT a, b, d, count() AS cnt FROM source

GROUP BY a, b, d;



Явная таблица dest



Неявная таблица `.inner.mv1`

- не работает optimize_move_to_prewhere при запросах к mv приходится писать from `.inner.mv1`
- paботает populate
- drop table mv1 неявный сторадж удаляется

Явная таблица dest

- не работает populate (можно сделать insert)
- drop table mv1 явный сторадж HE удаляется

• insert into mv1 / insert into `.inner.mv1` / insert into dest

Populate

• Ни разу не пригодился

• Таблица-сторадж не знает про MV и select в MV !!!

Неправильный ключ (order by) таблицы

Неправильно

```
CREATE MATERIALIZED VIEW mv1
ENGINE = SummingMergeTree
PARTITION BY toYYYYMM(d)
ORDER BY (a, b)
AS
SELECT a, b, d, count() AS cnt
FROM source
GROUP BY a, b, d;
```

Правильно

```
CREATE MATERIALIZED VIEW mv1
ENGINE = SummingMergeTree
PARTITION BY toYYYYMM(d)
ORDER BY (a, b, d)
AS
SELECT a, b, d, count() AS cnt
FROM source
GROUP BY a, b, d;
```

Двойные группировки (неправильно)

```
CREATE MATERIALIZED VIEW mv1
                                                  insert into source
ENGINE = AggregatingMergeTree
                                                  values (now()), (now())
PARTITION BY to YYYYMM (hour)
ORDER BY hour
                                                  результат \max by hour = 2
POPULATE
AS
SELECT toStartOfHour(time) hour,
       maxState(cnt by minute) max by hour,
                                                  insert into source
       sumState(cnt by minute) cnt by hour
                                                  values (now());
FROM
                                                  insert into source
                                                  values (now());
    SELECT minute, count() AS cnt by minute
    FROM source
                                                  результат \max by hour = 1
    GROUP BY minute
GROUP BY hour
```

MV никак не связан с исходной таблицей !!!

• MV не следит за Replacing/CollapsingMT в исходной таблице

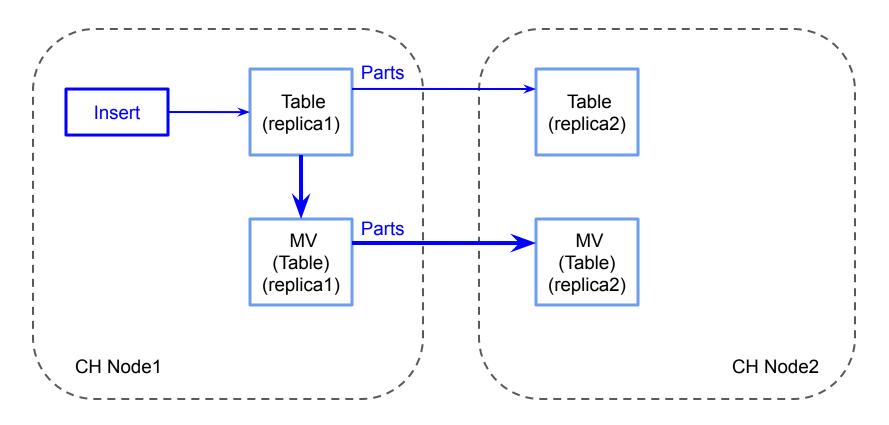
 MV ничего не знает про truncate / alter delete / alter update / drop partition / drop table / rename исходной таблицы

• MV может хранить данные за другой период

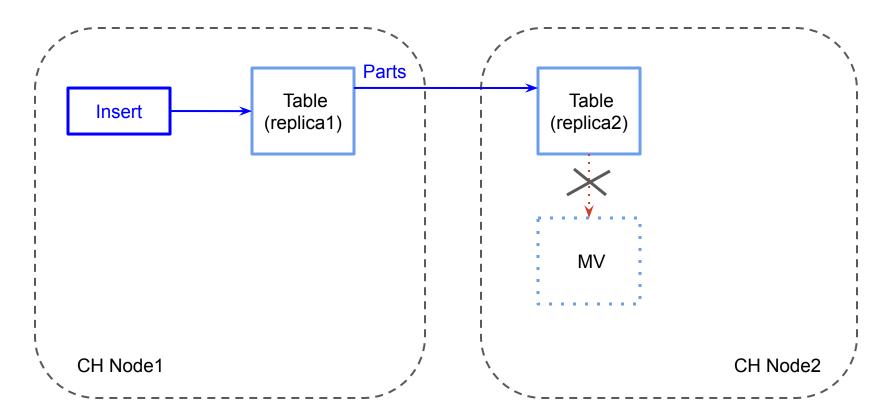
MV использует имена полей !!!

```
CREATE MATERIALIZED VIEW
                                         CREATE MATERIALIZED VIEW
mv1 (a Int64, d Date, cnt Int64)
                                         mv1 (a Int64, d Date, cnt Int64)
ENGINE = SummingMergeTree
                                         ENGINE = SummingMergeTree
PARTITION BY toYYYYMM(d)
                                         PARTITION BY to YYYYMM (d)
ORDER BY (a, d)
                                         ORDER BY (a, d)
POPULATE
                                         POPULATE
AS
                                         AS
SELECT a, d, count()
                                         SELECT a, d, count() as cnt
FROM source
                                         FROM source
GROUP BY a, d;
                                         GROUP BY a, d;
```

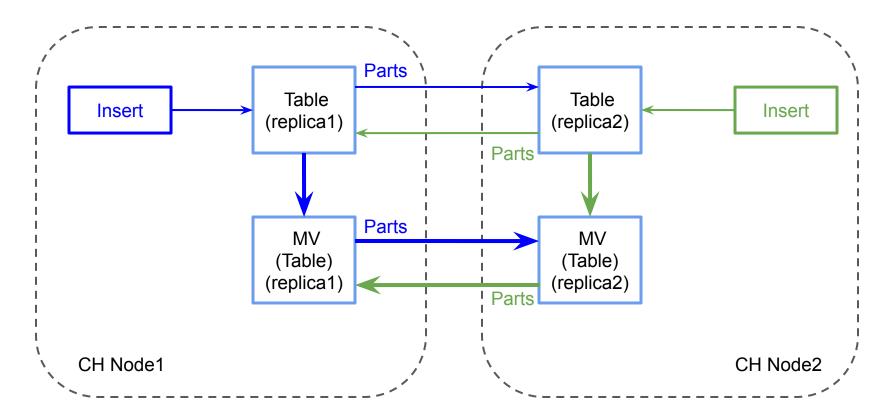
ReplicatedMT и MV



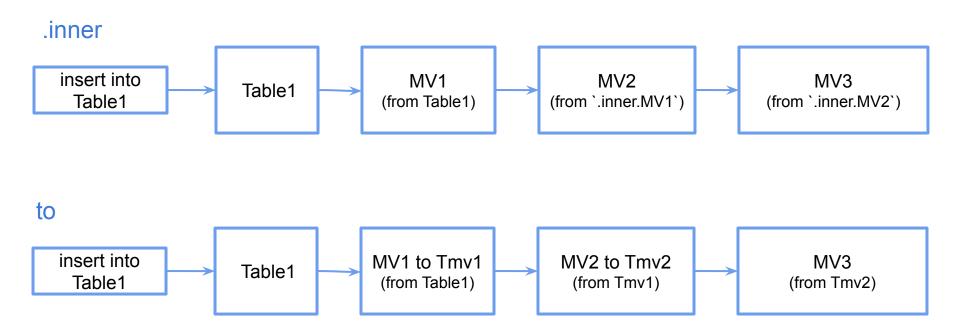
ReplicatedMT и MV



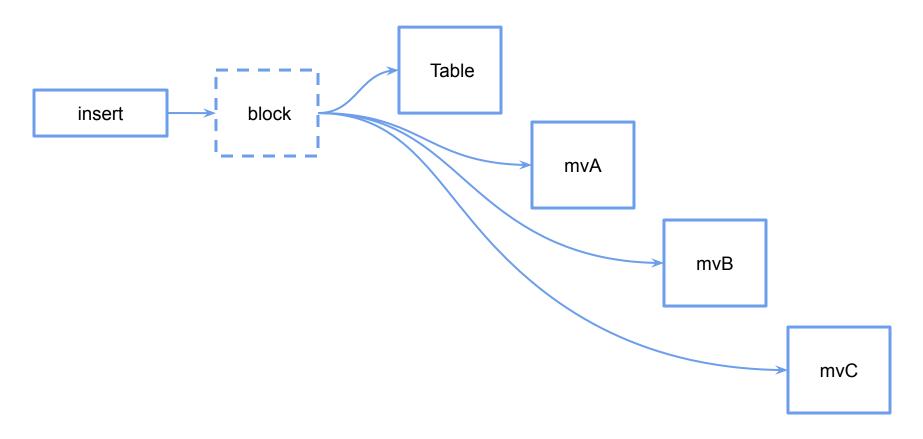
ReplicatedMT и MV



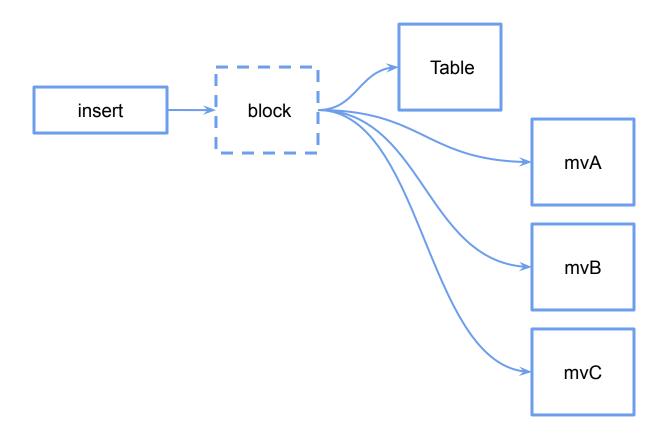
Каскадирование MV



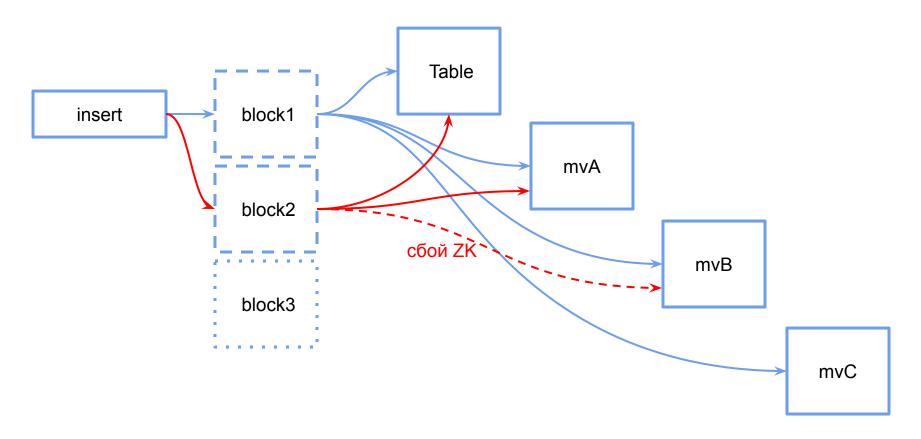
Порядок вставки в MV



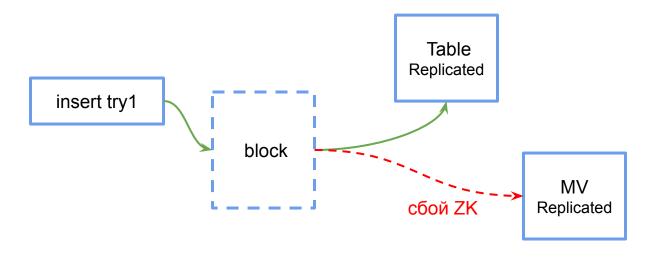
Порядок вставки в MV и parallel_view_processing



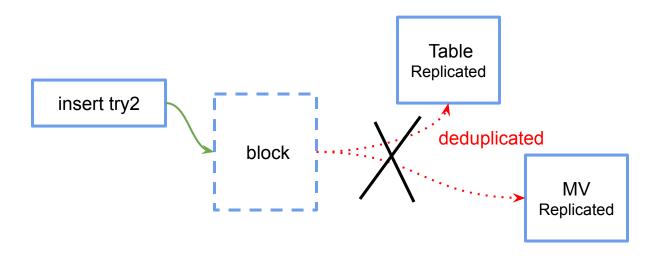
Порядок вставки в MV



MV insert deduplication



MV insert deduplication



napametp deduplicate_blocks_in_dependent_materialized_views = 0 (default)

JOIN B MV

- Несколько блоков -> несколько JOIN
- Внешние словари и dictGet или engine=Join и joinGet
- MV на обе таблицы в JOIN (2 MV to dest + RIGHT JOIN)

MV from distributed / To distributed

- From Local* to Local
- From Distributed to Distributed
- From Local* to Distributed
- From Distributed to Local

^{* (} то что Local за Distributed -- в этом случае несущественно)

Как изменить MV

Неявная таблица (.inner.mv1)

- 1. останавливаем insert-ы
- 2. detach table mv1
- 3. alter table `.inner.mv1`
- 4. attach materialized view mv1
- 4a. поправить .sql + attach table

Явная таблица (TO dest)

- 1. останавливаем insert-ы
- 2. alter table dest
- 3. drop table mv1
- 4. create materialized view mv1

Как изменить MV

```
Неявная таблица (.inner.mv1)
                                        Явная таблица (TO dest)
DETACH TABLE mv1
                                         ALTER TABLE dest
                                             ADD COLUMN b Int64 AFTER a,
ALTER TABLE `.inner.mv1`
                                             MODIFY ORDER BY (a, b)
    ADD COLUMN b Int64 AFTER a,
    MODIFY ORDER BY (a, b)
                                        DROP TABLE MV1
ATTACH MATERIALIZED VIEW mv1
                                        CREATE MATERIALIZED VIEW mv1
ENGINE = SummingMergeTree
                                         TO dest
ORDER BY (a, b) AS
                                         SELECT a, b, sum(amount) AS s
SELECT a, b, sum(amount) AS s
                                         FROM source
FROM source
                                         GROUP BY a, b
GROUP BY a, b
```

Как изменить MV

```
SET allow_experimental_alter_materialized_view_structure = 1;
ALTER TABLE mv1 MODIFY QUERY
    SELECT v * 2 as v, 1 as v2
    FROM source;
```

- создаем MV с where date >= в_будущем
 ждем в_будущем, вставляем все за до в_будущем
- останавливаем инсерты вставляем во вью все за сегодня запускаем инсерты вставляем все за до сегодня
- snapshot / alter table freeze
- создаем MV вставляем все за до вчера , (до начала партиции) ждем новый месяц (новую партицию) дропаем неполную партицию, пересчитываем

```
CREATE TABLE dest(a Int64, d Date, cnt UInt64)
ENGINE = SummingMergeTree
PARTITION BY to YYYYMM (d) ORDER BY (a, d);
-- создаем MV с where date >= в будущем
CREATE MATERIALIZED VIEW mv1 TO dest AS
SELECT a, d, count() AS cnt
FROM source
WHERE d \ge '2020-10-02'
GROUP BY a, d;
-- наступает 2020-10-02
INSERT INTO dest -- вставляем все за до в будущем
SELECT a, d, count() AS cnt
FROM source
WHERE d < '2020-10-02' -- кусками по 1-му месяцу (или .. дню)
GROUP BY a, d;
```

```
CREATE TABLE dest(a Int64, d Date, cnt UInt64)
ENGINE = SummingMergeTree
PARTITION BY to YYYYMM (d) ORDER BY (a, d);
-- останавливаем инсерты
INSERT INTO dest -- вставляем во вью все за сегодня
SELECT a, d, count() AS cnt
FROM source
WHERE d = today()
GROUP BY a, d;
CREATE MATERIALIZED VIEW mv1 TO dest AS
SELECT a, d, count() AS cnt
FROM source
GROUP BY a, d;
-- запускаем инсерты
INSERT INTO dest -- вставляем все за до сегодня
SELECT a, d, count() AS cnt
FROM source
WHERE d < today() -- кусками по 1-му месяцу (или .. дню)
GROUP BY a, d;
```

```
CREATE TABLE dest(a Int64, d Date, cnt UInt64)
ENGINE = SummingMergeTree
PARTITION BY to YYYYMM (d) ORDER BY (a, d);
-- останавливаем инсерты
ALTER TABLE dest FREEZE;
CREATE MATERIALIZED VIEW mv1 TO dest AS
SELECT a, d, count() AS cnt
FROM source
GROUP BY a, d;
-- запускаем инсерты
create table temp snapshot as source;
-- mv .../shadow/1/data/default/source/* .../data/default/temp snapshot/detached/
-- cd .../data/default/temp snapshot/detached/
-- let i=1; for f in `ls -1`; do echo $i $f;((i++)); echo "alter table temp snapshot attach
part '$f';"|clickhouse-client; done
INSERT INTO dest -- вставляем во выю из temp snapshot
SELECT a, d, count() AS cnt
FROM temp snapshot -- кусками по 1-му месяцу (или .. дню)
GROUP BY a, d;
```

```
CREATE TABLE dest(a Int64, d Date, cnt UInt64)
ENGINE = SummingMergeTree
PARTITION BY to YYYYMM(d) ORDER BY (a, d);
CREATE MATERIALIZED VIEW mv1 TO dest AS
SELECT a, d, count() AS cnt
FROM source
GROUP BY a, d;
INSERT INTO dest -- вставляем все за предыдущий период
SELECT a, d, count() AS cnt
FROM source
WHERE d < '2020-10-01' -- кусками по 1-му месяцу (или .. дню)
GROUP BY a, d;
-- наступает 2020-11-01
ALTER TABLE dest DROP PARTITION 202010;
INSERT INTO dest
SELECT a, d, count() AS cnt
FROM source
WHERE d \ge 2020-10-01' and d < 2020-11-01'
GROUP BY a, d;
```

полезняшки

Alter MV and SummingMergeTree

https://gist.github.com/den-crane/3a8d57253711e277b2a60a75b5dfeef6

How to convert not replicated MV to replicated

https://gist.github.com/den-crane/80cb95e74f046be2d8ffae58d9b04e8f

How to alter MV

how to convert MV with implicit storage .inner to explicit storage (with TO) https://gist.github.com/den-crane/431010ca08b9e51b960e55344b1dbbe3

MV_populating_with_freeze

https://gist.github.com/den-crane/64c05be67ef8a926242011188ce96f44

Several_MV_one_internal_store

https://gist.github.com/den-crane/005633470c70877dd28c00211cd9fcfb

Populate AggregatingMergeTree through null table

https://gist.github.com/den-crane/f7382cd4f1f859ff6ac46afe7dc9925a

Rename MV in CH before 19.8.3.8

https://gist.github.com/den-crane/d9a0bceb58a7d5314ea66e8edbebc9f7

